

```

import pandas as pd
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.tree import DecisionTreeRegressor as dt
from sklearn.ensemble import RandomForestRegressor as rf
from sklearn.metrics import r2_score
import warnings
warnings.filterwarnings("ignore")
import matplotlib.pyplot as plt
import numpy as np

```

```

z = pd.read_csv(r"C:\Users\skj_h\OneDrive\Desktop\dataset\
boston_housing.csv")
z

```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD
TAX \									
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0
296.0									
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0
242.0									
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0
242.0									
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0
222.0									
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0
222.0									
..
.									
501	0.06263	0.0	11.93	0.0	0.573	6.593	69.1	2.4786	1.0
273.0									
502	0.04527	0.0	11.93	0.0	0.573	6.120	76.7	2.2875	1.0
273.0									
503	0.06076	0.0	11.93	0.0	0.573	6.976	91.0	2.1675	1.0
273.0									
504	0.10959	0.0	11.93	0.0	0.573	6.794	89.3	2.3889	1.0
273.0									
505	0.04741	0.0	11.93	0.0	0.573	6.030	80.8	2.5050	1.0
273.0									
	PTRATIO	B	LSTAT	MEDV					
0	15.3	396.90	4.98	24.0					
1	17.8	396.90	9.14	21.6					
2	17.8	392.83	4.03	34.7					
3	18.7	394.63	2.94	33.4					
4	18.7	396.90	5.33	36.2					
..					
501	21.0	391.99	9.67	22.4					
502	21.0	396.90	9.08	20.6					

```
503      21.0  396.90   5.64  23.9
504      21.0  393.45   6.48  22.0
505      21.0  396.90   7.88  11.9
```

```
[506 rows x 14 columns]
```

```
z.isnull().sum()
```

```
CRIM      0
ZN        0
INDUS     0
CHAS      0
NOX       0
RM        0
AGE       0
DIS       0
RAD       0
TAX       0
PTRATIO   0
B         0
LSTAT     0
MEDV      0
dtype: int64
```

```
z.shape
```

```
(506, 14)
```

```
z.size
```

```
7084
```

```
z.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 506 entries, 0 to 505
```

```
Data columns (total 14 columns):
```

#	Column	Non-Null Count	Dtype
0	CRIM	506 non-null	float64
1	ZN	506 non-null	float64
2	INDUS	506 non-null	float64
3	CHAS	506 non-null	float64
4	NOX	506 non-null	float64
5	RM	506 non-null	float64
6	AGE	506 non-null	float64
7	DIS	506 non-null	float64
8	RAD	506 non-null	float64
9	TAX	506 non-null	float64
10	PTRATIO	506 non-null	float64
11	B	506 non-null	float64

```

12  LSTAT      506 non-null    float64
13  MEDV       506 non-null    float64
dtypes: float64(14)
memory usage: 55.5 KB

z.ndim
2


z.dtypes
CRIM      float64
ZN        float64
INDUS     float64
CHAS      float64
NOX       float64
RM        float64
AGE       float64
DIS       float64
RAD       float64
TAX       float64
PTRATIO   float64
B         float64
LSTAT     float64
MEDV      float64
dtype: object

abs(z.corr()["MEDV"]).sort_values(ascending = False)
MEDV      1.000000
LSTAT     0.737663
RM        0.695360
PTRATIO   0.507787
INDUS     0.483725
TAX       0.468536
NOX       0.427321
CRIM      0.388305
RAD       0.381626
AGE       0.376955
ZN        0.360445
B         0.333461
DIS       0.249929
CHAS      0.175260
Name: MEDV, dtype: float64

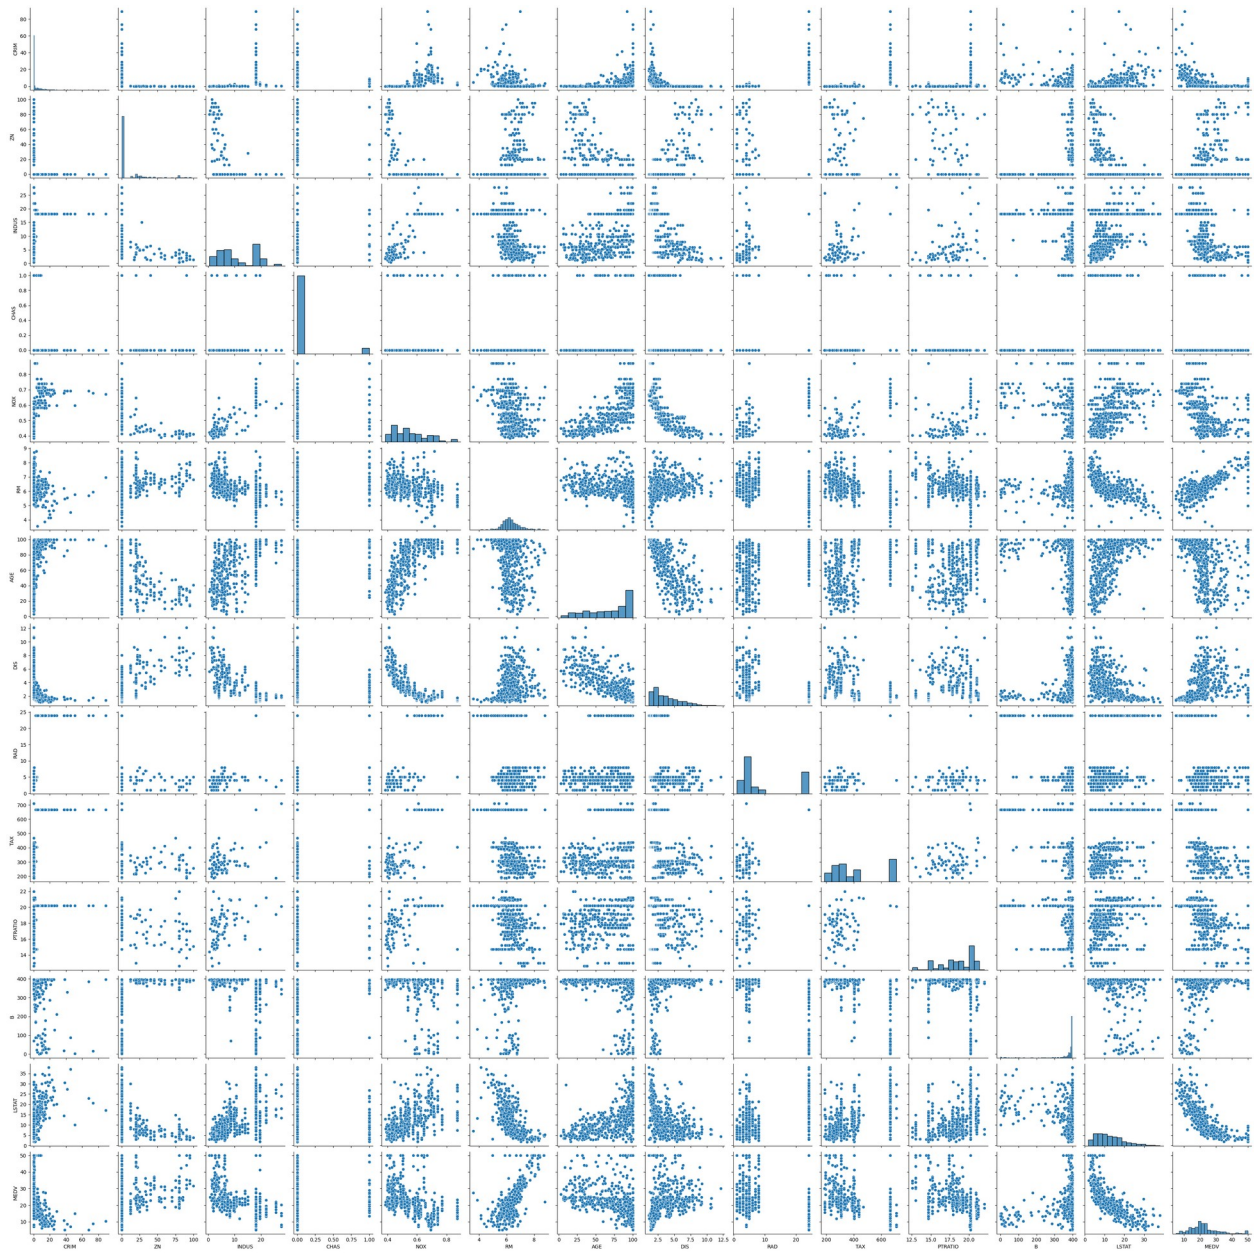
sns.scatterplot(x = z["LSTAT"], y = z["MEDV"], data = z)

<Axes: xlabel='LSTAT', ylabel='MEDV'>

```




```
sns.pairplot(z)  
<seaborn.axisgrid.PairGrid at 0x26bb980fc80>
```



```
plt.figure(figsize = (10, 10))
sns.heatmap(z.corr(), annot = True, alpha = 1)
```

<Axes: >




```
z.columns
```

```
Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS',  
      'RAD', 'TAX',  
      'PTRATIO', 'B', 'LSTAT', 'MEDV'],  
      dtype='object')
```

```
sns.regplot(x = z["LSTAT"], y = z["MEDV"], data = z, line_kws =  
            {"color" : "red"})
```

```
<Axes: xlabel='LSTAT', ylabel='MEDV'>
```



```
z.columns
Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS',
      'RAD', 'TAX',
      'PTRATIO', 'B', 'LSTAT', 'MEDV'],
      dtype='object')

x = z[["LSTAT", "MEDV"]]
X = x
Y = x["MEDV"]

x_train, x_test, y_train, y_test = train_test_split(X, Y, train_size =
0.7, test_size = 0.3, random_state = 100)

x_train.shape
(354, 2)

x_train = x_train.drop(["MEDV"], axis = 1)
x_test = x_test.drop(["MEDV"], axis = 1)

y_train = np.array(y_train).reshape(-1, 1)
y_test = np.array(y_test).reshape(-1, 1)
```

```

n = LinearRegression()
n.fit(x_train, y_train)

LinearRegression()

y_predict_train = n.predict(x_train)
r2_train__LinearRegression = r2_score(y_true = y_train, y_pred =
y_predict_train)

y_predict_test = n.predict(x_test)
r2_test__LinearRegression = r2_score(y_true = y_test, y_pred =
y_predict_test)

a = dt()
a.fit(x_train, y_train)

DecisionTreeRegressor()

y_predict_train1 = a.predict(x_train)
r2_train__DecisionTree = r2_score(y_true = y_train, y_pred =
y_predict_train1)

y_predict_test1 = a.predict(x_test)
r2_test__DecisionTree = r2_score(y_true = y_test, y_pred =
y_predict_test1)

b = rf()
b.fit(x_train, y_train)

RandomForestRegressor()

y_predict_train2 = b.predict(x_train)
r2_train__RandomForest = r2_score(y_true = y_train, y_pred =
y_predict_train2)

y_predict_test2 = b.predict(x_test)
r2_test__RandomForest = r2_score(y_true = y_test, y_pred =
y_predict_test2)

z5 = {"Model name" : ["LinearRegression", "Decision Tree Regression",
"Random Forest Regression"],
      "Evaluation on training dataset" : [r2_train__LinearRegression,
r2_train__DecisionTree,
r2_train__RandomForest],
      "Evaluation on testing dataset" : [r2_test__LinearRegression,
r2_test__DecisionTree,
r2_test__RandomForest]
      }

List_model_with_r2_score = pd.DataFrame(z5)
List_model_with_r2_score

```


	Model name	Evaluation on training dataset	\
0	LinearRegression	0.546342	
1	Decision Tree Regression	0.960354	
2	Random Forest Regression	0.910186	

	Evaluation on testing dataset
0	0.532230
1	0.471643
2	0.565357